

---

# *Data Stream Algorithms I*

---

Joaquim Madeira

Version 0.2 – November 2018

---

# Overview

- The data stream model
- Finding frequent items
- The MAJORITY problem
- The FREQUENT problem

# Data Streams

- Many data generation processes can be modeled as **data streams**
  - Huge numbers of **simple** pieces of data
  - Arriving at **enormous rates**
  - Taken together lead to **a complex whole**
- Hundreds of gigabytes per day or higher !

# Data Streams

- Sequence of **queries** posed to an Internet search engine
- Collection of **transactions** across all branches of a supermarket chain
- Sequence of **packets** in network traffic monitoring
- ...

# Data Streams

- Such data may be archived and indexed within a **data warehouse**
- BUT, it may also be important to process it **“as it happens”**
- Up to the minute **analysis** and **statistics** on current **trends**

# Data Streams

- **Quick response** to each new piece of information
- **Resources** used **very small** when compared to the total quantity of data

# The streaming model

- Data arrives in a **streaming** fashion
  - Scan the sequence in the given order
  - **No random access** to the data tokens !
- Must be **processed on the fly** !
- **Accurate** computations

# The streaming model

- Compute some function  $\Phi(\sigma)$  of a **massively long** input stream  $\sigma$
- Make just **one pass** over  $\sigma$  !
- Goal:
  - Use resources ( **space** and **time** ) **sublinear** on the size of the input !



# The streaming model

- When to produce **output** ?
- At the **end** of the stream
- When **queried** on the stream **prefix** observed so far
- Whenever there is a stream **update**
- On a “**sliding window**” of the most recent updates

# The basic streaming model

- The data stream:

$$\sigma = \langle a_1, a_2, \dots, a_m \rangle$$

- Each data token  $a_i$  is drawn from a set of  $n$  elements
- Goal:
  - Process  $\sigma$  using a small amount of memory  $s$
  - I.e., make  $s$  much smaller than  $m$  and  $n$  !

# The quality of an algorithm's answer

- $\Phi(\sigma)$  is usually a **real-valued** function
- Allow for
  - Computing an **estimate** or **approximation** of  $\Phi(\sigma)$
  - Possibly using randomized algorithms
    - That may err with a small, but controllable probability
- How to evaluate the **quality** of the result ?

# The quality of an algorithm's answer

- $A(\sigma)$  is the output of a randomized algorithm
  - It is a random variable !

- $(\varepsilon, \delta)$ -approximation of  $\Phi(\sigma)$

$$P\left(\left|\frac{A(\sigma)}{\Phi(\sigma)} - 1\right| > \varepsilon\right) \leq \delta$$

- $(\varepsilon, \delta)$ -additive-approximation of  $\Phi(\sigma)$


$$P(|A(\sigma) - \Phi(\sigma)| > \varepsilon) \leq \delta$$

# Finding frequent items

- The frequent items / “heavy-hitters” problem
- Given a sequence of items, identify those which occur most frequently
- More formally :
- Find all items whose frequency exceeds a specified fraction of the total number of items

# Finding frequent items

**Figure 1. A stream of items defines a frequency distribution over items. In this example, with a threshold of  $\phi = 20\%$  over the 19 items grouped in bins, the problem is to find all items with frequency at least 3.8—in this case, the green and red items (middle two bins).**



[Cormode and Hadjieleftheriou]

# Finding frequent items

- Network packet monitoring
  - Frequent items represent the **heaviest bandwidth users**
- Queries made to a search engine
  - Frequent items are the currently **popular terms**
- ...

# Finding frequent items

- Counter-based algorithms
  - Track and maintain counts associated with a (varying) subset of stream items
- Sketch algorithms
  - Randomized approach
  - Do not explicitly store stream elements
- Other approaches



# Finding frequent items

- Given a stream:  $\sigma = \langle a_1, a_2, \dots, a_m \rangle$
- It induces a frequency vector:

$$\mathbf{f} = (f_1, f_2, \dots, f_n)$$

$$f_1 + f_2 + \dots + f_n = m$$

- The MAJORITY problem
  - If  $\exists j : f_j > \frac{m}{2}$ , then output  $j$ , otherwise output null
- The FREQUENT problem, with parameter  $k$

# The MAJORITY problem

- Applications ?
- Elections
- Fault-tolerant computing
  - Perform multiple redundant computations
  - Check if a majority of the results agree
- ...

# The MAJORITY problem

- Naïve algorithm for a non-sorted list of values
- Sort the list
- If there is a majority value, it is now the middle value
  - Odd vs even number of list elements
- $O(n \log n)$
- BUT, not useful for data streams !

# The MAJORITY problem

- Naïve algorithm
- $n$  frequency counters
- Three-step algorithm
  - Scan the sequence and increment the counters
  - Scan the counters and find the most frequent element
  - Check if it is the majority element :  $> ( m / 2 )$
- Efficiency ?

# The MAJORITY problem

- Boyer & Moore : A fast majority vote alg.
  - 1980
  - <http://www.cs.utexas.edu/~moore/best-ideas/mjrty/>
- **Provided there is such an element**, it decides which sequence element is in the majority
- **Two-pass** algorithm
  - Scan the sequence to identify the **majority candidate**
  - Scan, **again**, the sequence, to verify if that candidate is indeed in the majority

# MJRTY – A fast majority vote alg.


## // Initialization

```
candidate = null; counter = 0;
```

## // First-pass

```
while ( not end of sequence )  
    x = current_token();  
    if ( counter == 0 )  
        then candidate = x; counter = 1;  
    else    if ( candidate == x )  
        then counter++;  
        else counter--;
```

# MJRTY – A fast majority vote alg.



[Wikipedia]

# MJRTY – A fast majority vote alg.

// Second-pass

```
counter = 0;
```

```
while ( not end of sequence and counter < (1 + m / 2) )
```

```
    x = current_token();
```

```
    if ( candidate == x )
```

```
        then counter++;
```

■ Efficiency ?



# MJRTY – A fast majority vote alg.

- Can we skip the second pass?
  - Find a **counter-example** !
- $O(1)$  extra space
- $O(n)$  time
- **BUT**, we cannot perform a second pass over a data stream...
  - However, we have a “**partial guarantee**”

# Tasks – The MAJORITY problem

- Implement the naïve algorithm
- Implement the Boyer & Moore algorithm
- Compare their results and running times
  - For random strings over a given alphabet
- For the B & M algorithm, check how many times the majority candidate was indeed the majority

# The FREQUENT problem

- The FREQUENT problem, with parameter  $k$ 
  - Output the set  $\{j : f_j > m/k\}$
- It solves the MAJORITY problem !
- Similar naïve algorithm !
- Can we do better ?

# Frequency estimation


- The FREQUENCY-ESTIMATION problem
  - Process the stream  $\sigma$
  - Establish an **estimate** for the frequency of any stream token
- Misra & Gries : Finding repeated elements
  - **1982**
  - <http://www.sciencedirect.com/science/article/pii/0167642382900120>
- **One-pass** algorithm

# The Misra & Gries algorithm

- **Parameter  $k$**  controls the quality of the results given
- It maintains an **associative array**
  - The **keys** are **tokens** seen in the stream
  - Array **values** are **counters** associated with the keys / tokens
- At most  **$(k - 1)$  counters**, at any time

# The Misra & Gries algorithm

**Figure 2. Counter-based data structure:** the blue (top) item is already stored, so its count is incremented when it is seen. The green (middle) item takes up an unused counter, then a second occurrence increments it.



[Cormode and  
Hadjieleftheriou]

# The Misra & Gries algorithm

---

**Algorithm 1:** FREQUENT( $k$ )

---

```
 $n \leftarrow 0;$   
 $T \leftarrow \emptyset;$   
foreach  $i$  do  
   $n \leftarrow n + 1;$   
  if  $i \in T$  then  
     $c_i \leftarrow c_i + 1;$   
  else if  $|T| < k-1$  then  
     $T \leftarrow T \cup \{i\};$   
     $c_i \leftarrow 1;$   
  else forall  $j \in T$  do  
     $c_j \leftarrow c_j - 1;$   
    if  $c_j = 0$  then  $T \leftarrow T \setminus \{j\};$ 
```

---

[Cormode and  
Hadjieleftheriou]

# The Misra & Gries algorithm

## // Initialization

A = empty associative array;

## // Processing

while ( not end of sequence )

    j = current\_token();

    if ( j in keys(A) ) then A[ j ] = A[ j ] + 1;

    else if ( | keys(A) | < ( k - 1 ) ) then A[ j ] = 1;

        else for each i in keys(A) do

            A[ i ] = A[ i ] - 1;

            if ( A[ i ] == 0 ) then remove i from A;

## // Output

if( a in keys(A) ) then freq\_estimate = A[ a ];

else freq\_estimate = 0;



# The Misra & Gries algorithm

- The algorithm, with parameter  $k$ , provides, for any token  $j$ , a freq. estimate  $f_j^*$  satisfying

$$f_j - \frac{m}{k} \leq f_j^* \leq f_j$$

- If some token has  $f_j > m / k$ , its counter  $A[j]$  will be positive
- With an additional pass through the stream, to count the exact frequencies, we can now solve the FREQUENT problem !

# Tasks – The Misra & Gries algorithm

- Implement the naïve algorithm
- Implement the Misra & Gries algorithm
  - Choose an appropriate data structure for the associative array
- Test them, using the same input sequences as for the B & M algorithm
  - For different  $k$  values !!

# Misra & Gries – Recap

- Finds up to  $(k - 1)$  items that occur more than a  $1/k$  fraction of the time in the input
- Keeps, at most,  $(k - 1)$  candidates at the same time
- No item with frequency  $m / k$  is missed
- Algorithm “rediscovered” twice in 2002 !

# Implementation issues

## ■ Basic steps

- ❑ Lookup for an item
- ❑ Update a counter
- ❑ Decrement all counters
- ❑ Delete an item with zero counts

## ■ How to ?

- ❑ Optimize speed and space

# Implementation issues – Lookup

- Which **dictionary** data structure ?
- Misra & Gries used a **balanced search tree**
  - Worst and average case are  **$O(\log k)$**
- **Hash table** : hash to  **$O(k)$**  buckets
  - **Collisions / deletions** : how to handle ?
  - Use **chaining** ?
  - Optimizations ?
- Other ?

# Implementation issues – Decrement

- Iterate through all counters :  $O(k)$
- BUT, it happens  $O(n/k)$  times
- Optimize ?
- Use a linked **list of lists** to keep elements grouped by their **frequency counts**
- Memory space **overhead**
  - ❑ Circular linked lists
  - ❑ Also, pointers **to and from** hash table

# Additional algorithms

- There are other algorithms which can be regarded as variations of Misra & Gries' algorithm :
- Lossy-Counting
  - Manku and Motwani, 2002
- Space-Saving
  - Metwally et al., 2005

# The Manku & Motwani algorithm

---

**Algorithm 2:** LOSSYCOUNTING( $k$ )

---

```
 $n \leftarrow 0; \Delta \leftarrow 0; T \leftarrow \emptyset;$   
foreach  $i$  do  
   $n \leftarrow n + 1;$   
  if  $i \in T$  then  $c_i \leftarrow c_i + 1;$   
  else  
     $T \leftarrow T \cup [i];$   
     $c_i \leftarrow 1 + \Delta;$   
  
  if  $\lfloor n/K \rfloor \neq \Delta$  then  
     $\Delta \leftarrow \lfloor n/k \rfloor;$   
    forall  $j \in T$  do  
      if  $c_j < \Delta$  then  $T \leftarrow T \setminus [j];$ 
```

---

[Cormode and  
Hadjieleftheriou]



# Manku & Motwani – Lossy-Counting

- Keep item names and counts
  - Counter value is a **lower bound** – initially **zero**
- And an “implicit” **delta** value
- A **new item** – what to do ?
- If it has a counter, **increment** counter
- Otherwise, **initialize** with a count of **1 + delta**
- Whenever **delta increases** :
  - **Delete** tuples with a count **smaller** than delta

# Manku & Motwani – Lossy-Counting

- Deleting tuples reduces the **required space** !
- Monitored items can have their frequencies **overestimated** by no more than  $n / k = \epsilon \times n$
- BUT **never** underestimated !!

# The Metwally et al. algorithm

---

**Algorithm 3:** SPACEAVING( $k$ )

---

```
 $n \leftarrow 0;$   
 $T \leftarrow \emptyset;$   
foreach  $i$  do  
     $n \leftarrow n + 1;$   
    if  $i \in T$  then  $c_i \leftarrow c_i + 1;$   
    else if  $|T| < k$  then  
         $T \leftarrow T \cup \{i\};$   
         $c_i \leftarrow 1;$   
    else  
         $j \leftarrow \arg \min_{j \in T} c_j;$   
         $c_i \leftarrow c_j + 1;$   
         $T \leftarrow T \cup \{i\} \setminus \{j\};$ 
```

---

[Cormode and  
Hadjieleftheriou]

# Metwally et al. – Space-Saving

- Keep  $k=1/\epsilon$  item names and counts
  - Initially zero
- Count first  $k$  items exactly !
- A new item – what to do ?
- If it has a counter, increment counter
- Otherwise, replace item with least count
- And increment count

# Metwally et al. – Space-Saving

- Counters **sum** to  **$n$**  !
- Average count value is  **$n / k = \epsilon \times n$** 
  - Smallest count **min** cannot be larger than  **$\epsilon \times n$**
- True count of an uncounted item is between 0 and  **$\epsilon \times n$**
- All items whose true count is  **$> \epsilon \times n$**  are **stored** !

# Tasks

- Implement the **Lossy-Counting** and the **Space-Saving** algorithms
- Test them, using the same input sequences as for the M & G algorithm
- Compare the behavior of the three algorithms

# Implementation issues

- Similar to Misra & Gries
- Finding the min item is a standard problem
  - Use a **min-heap** !
  - Binary, binomial, Fibonacci, ... ?
  - **$O(\log k)$**

# Question

- What can you say about the **estimated counts** for **items** which are **stored** by the algorithms **early** in the stream and are **not removed** ?



# Question

- Have we been discussing **deterministic** algorithms or **randomized/probabilistic** algorithms ?

# Experimental comparison

- Cormode & Hadjieleftheriou
  - VLDB 2008 - <https://dl.acm.org/citation.cfm?id=1454225>
  - CACM 2009 - <https://dl.acm.org/citation.cfm?id=1562789>
- **SPACESAVING** has benefits over others !
- Very fast: 20M – 30M updates *per second*
- Implementation choices: **speed** vs **space**
  - E.g., a heap or lists of items grouped by frequencies

# 2017 – Recent progress

## A High-Performance Algorithm for Identifying Frequent Items in Data Streams

Daniel Anderson  
Georgetown University

Pryce Bevan  
Georgetown University

Kevin Lang  
Oath Research

Edo Liberty\*  
Amazon

Lee Rhodes  
Oath

Justin Thaler  
Georgetown University

### ABSTRACT

Estimating frequencies of items over data streams is a common building block in streaming data measurement and analysis. Misra and Gries introduced their seminal algorithm for the problem in 1982, and the problem has since been revisited many times due its practicality and applicability. We describe a highly optimized version of Misra and Gries' algorithm that is suitable for deployment in industrial settings. Our code is made public via an open source library called Data Sketches that is already used by several companies and production systems.

been studied intensely [6, 7, 9, 13, 14, 17, 21, 31–35]. These algorithms process a massive dataset in a single pass, and compute very small *summaries* of the dataset, from which it is possible to derive accurate—though approximate—answers to frequent items queries and point queries.

It may seem as though streaming frequency approximation is well-understood, with little room for further insight or improvement. However, when we set about implementing an algorithm suitable for industrial use on web-scale data, we found that existing algorithms have two significant shortcomings. First, they are not

---

\*Research performed while at Yahoo Research.

<https://dl.acm.org/citation.cfm?doid=3131365.3131407>

<https://datasketches.github.io/>

# References

- R. Boyer & J. Moore, MJRTY – A fast majority vote algorithm, in *Automated Reasoning: Essays in Honor of Woody Bledsoe*, Springer, 1991
- J. Misra & D. Gries, Finding repeated elements, *Science of Computer Programming*, Vol. 2, 1982
- G. Cormode & M. Hadjieleftheriou, Finding the frequent items in streams of data, *Commun. ACM*, Vol. 52, N. 10, 2009